
Tranalyzer2

kafkaSink



Apache Kafka



Tranalyzer Development Team

Contents

1	kafkaSink	1
1.1	Description	1
1.2	Dependencies	1
1.3	Services Initialization	1
1.4	Configuration Flags	2
1.5	Plugin Report Output	2
1.6	Example	2

1 kafkaSink

1.1 Description

The kafkaSink plugin outputs flows to an Apache Kafka event streaming platform.

1.2 Dependencies

1.2.1 External Libraries

This plugin depends on the **librdkafka** library.

OPT2=1		
Ubuntu:	sudo apt-get install	librdkafka-dev
Arch:	sudo pacman -S	librdkafka
Gentoo:	sudo emerge	librdkafka
openSUSE:	sudo zypper install	librdkafka-devel
Red Hat/Fedora¹:	sudo dnf install	librdkafka-devel
macOS²:	brew install	librdkafka

1.2.2 Core Configuration

This plugin requires the following core configuration:

- `$T2HOME/tranalyzer2/src/tranalyzer.h:`
 - `BLOCK_BUF=0`

1.3 Services Initialization

The kafkaSink plugin requires a ZooKeeper and a Kafka broker service running on `KAFKA_BROKERS` (default address and port are `127.0.0.1:9092`).

```
# Start the ZooKeeper server and send it to the background
$ zookeeper-server-start.sh /etc/kafka/zookeeper.properties &

# Start the Kafka server and send it to the background
$ kafka-server-start.sh /etc/kafka/server.properties &
```

¹If the `dnf` command could not be found, try with `yum` instead

²Brew is a packet manager for macOS that can be found here: <https://brew.sh>

1.4 Configuration Flags

The following flags can be used to control the output of the plugin:

Name	Default	Description
KAFKA_BROKERS	"127.0.0.1:9092"	Broker address(es) (comma separated list of host[:port])
KAFKA_TOPIC	"tranalyzer.flows"	Topic to produce to
KAFKA_PARTITION	-1	Target partition: ≥ 0: fixed partition -1: automatic partitioning (unassigned)
KAFKA_RETRIES	3	Max. number of retries when message production failed [0 - 255]
KAFKA_DEBUG	0	Print debug messages

1.4.1 Environment Variable Configuration Flags

The following configuration flags can also be configured with environment variables (ENVCTRL>0):

- KAFKA_BROKERS
- KAFKA_TOPIC
- KAFKA_PARTITION

1.5 Plugin Report Output

The following information is reported:

- Number of flows discarded

1.6 Example

In this example, the flows will be sent to Kafka to the `tranalyzer.flows` topic. In addition, Tranalyzer information ([INF]) and warnings ([WRN]) will be sent to the `tranalyzer.out`, while errors ([ERR]) will use the `tranalyzer.err` topic.

First, we want to prevent Tranalyzer from coloring the output:

```
$ t2conf tranalyzer2 -D T2_LOG_COLOR=0
```

In this example, only the `basicFlow` and `kafkaSink` plugins will be used:

```
$ t2build tranalyzer2 basicFlow kafkaSink
```

Now, the fun part! Run Tranalyzer as per usual, but redirect `stdout` and `stderr` to a `kcat`³ process, which will send the data to Kafka:

³kcat was formerly known as `kafkacat`

```
$ t2 -r file.pcap \  
  1> >(grep -F -e "[INF]" -e "[WRN]" | kcat -P -b 127.0.0.1:9092 -t tranalyzer.out) \  
  2> >(kcat -P -b 127.0.0.1:9092 -t tranalyzer.err)
```

The messages can now be consumed with the help of `kafka-console-consumer`.

```
# Consume messages for tranalyzer.flows topic
```

```
$ kafka-console-consumer \  
  --bootstrap-server localhost:9092 \  
  --from-beginning \  
  --topic tranalyzer.flows
```

```
# Consume messages for tranalyzer.out topic
```

```
$ kafka-console-consumer \  
  --bootstrap-server localhost:9092 \  
  --from-beginning \  
  --topic tranalyzer.out
```

```
# Consume messages for tranalyzer.err topic
```

```
$ kafka-console-consumer \  
  --bootstrap-server localhost:9092 \  
  --from-beginning \  
  --topic tranalyzer.err
```