

---


# Tranalyzer2

psqlSink



PostgreSQL

---



Tranalyzer Development Team

Contents

<b>1</b>	<b>psqlSink</b>	<b>1</b>
1.1	Description . . . . .	1
1.2	Dependencies . . . . .	1
1.3	Database Initialization . . . . .	1
1.4	Configuration Flags . . . . .	1
1.5	Insertion of Selected Fields Only . . . . .	2
1.6	Post-Processing . . . . .	3
1.7	Example . . . . .	3

# 1 psqlSink

## 1.1 Description

The psqlSink plugin outputs flows to a PostgreSQL database.

## 1.2 Dependencies

### 1.2.1 External Libraries

This plugin depends on the **libpq** library.

<b>Ubuntu:</b>	sudo apt-get install	libpq-dev
<b>Arch:</b>	sudo pacman -S	postgresql-libs
<b>Gentoo:</b>	sudo emerge	postgresql
<b>openSUSE:</b>	sudo zypper install	postgresql-devel
<b>Red Hat/Fedora<sup>1</sup>:</b>	sudo dnf install	libpq-devel
<b>macOS<sup>2</sup>:</b>	brew install	postgresql

### 1.2.2 Core Configuration

This plugin requires the following core configuration:

- `$T2HOME/tranalyzer2/src/tranalyzer:h:`
  - `BLOCK_BUF=0`

## 1.3 Database Initialization

The psqlSink plugin requires a PostgreSQL server running on `PSQL_HOST`, e.g., `127.0.0.1` on port `PSQL_PORT`, e.g., `5432`. In addition, a user with name `PSQL_USER`, e.g., `postgres`, and password `PSQL_PASS`, e.g., `postgres` **MUST** exist and be allowed to create databases. This can be achieved with the following commands:

**Linux:** `$ sudo -u postgres psql` **macOS:** `$ psql postgres`

```
postgres=# CREATE ROLE postgres WITH LOGIN PASSWORD 'postgres';
CREATE ROLE
postgres=# ALTER ROLE postgres CREATEDB;
ALTER ROLE
```

## 1.4 Configuration Flags

The following flags can be used to control the output of the plugin:

Name	Default	Description
<code>PSQL_OVERWRITE_DB</code>	<code>2</code>	<code>0</code> : abort if DB already exists

---

<sup>1</sup>If the `dnf` command could not be found, try with `yum` instead

<sup>2</sup>Brew is a packet manager for macOS that can be found here: <https://brew.sh>

Name	Default	Description
PSQL_OVERWRITE_TABLE	2	1: overwrite DB if it already exists 2: reuse DB if it already exists 0: abort if table already exists
PSQL_TRANSACTION_NFLOWS	40000	1: overwrite table if it already exists 2: append to table if it already exists 0: one transaction > 0: one transaction every $n$ flows
PSQL_QRY_LEN	32768	Max length for query
PSQL_HOST	"127.0.0.1"	Address of the database
PSQL_PORT	5432	Port of the database
PSQL_USER	"postgres"	Username to connect to DB
PSQL_PASS	"postgres"	Password to connect to DB
PSQL_DBNAME	"tranalyzer"	Name of the database
PSQL_TABLE_NAME	"flow"	Name of the table
PSQL_SELECT	0	Only insert specific fields into the DB
PSQL_SELECT_FILE	"psql-columns.txt"	Filename of the field selector (one column name per line)

#### 1.4.1 Environment Variable Configuration Flags

The following configuration flags can also be configured with environment variables (ENVCTRL>0):

- PSQL\_HOST
- PSQL\_PORT
- PSQL\_USER
- PSQL\_PASS
- PSQL\_DBNAME
- PSQL\_TABLE\_NAME
- PSQL\_SELECT\_FILE (require PSQL\_SELECT=1)

## 1.5 Insertion of Selected Fields Only

When PSQL\_SELECT=1, the columns to insert into the DB can be customized with the help of PSQL\_SELECT\_FILE. The filename defaults to psql-columns.txt in the user plugin folder, e.g., *~/tranalyzer/plugins*. The format of the file is simply one field name per line with lines starting with a '#' being ignored. For example, to only insert source and destination addresses and ports, create the following file:

```
# Lines starting with a '#' are ignored and can be used to add comments
srcIP
srcPort
dstIP
dstPort
```

## 1.6 Post-Processing

The following queries can be used to analyze bitfields in PostgreSQL:

- Select all A flows:  

```
SELECT to_hex("flowStat"::bigint), *  
FROM flow  
WHERE ("flowStat"::bigint & 1) = 0::bigint
```
- Select all IPv4 flows:  

```
SELECT *  
FROM flow  
WHERE ("flowStat"::bigint & x'4000'::bigint) != 0::bigint
```
- Select all IPv6 flows:  

```
SELECT to_hex("flowStat"::bigint), *  
FROM flow  
WHERE ("flowStat"::bigint & x'8000'::bigint) != 0::bigint
```

## 1.7 Example

```
# Run Tranalyzer  
$ t2 -r file.pcap  
  
# Connect to the PostgreSQL database  
$ psql -U postgres -d tranalyzer  
  
# Number of flows  
tranalyzer=# SELECT COUNT(*) FROM flow;  
  
# 10 first srcIP/dstIP pairs  
tranalyzer=# SELECT "srcIP", "dstIP" FROM flow LIMIT 10;  
  
# All flows from 1.2.3.4 to 1.2.3.5  
tranalyzer=# SELECT * FROM flow WHERE "srcIP" = '1.2.3.4' AND "dstIP" = '1.2.3.5';
```

For examples of more complex queries, have a look in \$T2HOME/scripts/t2fm/psql/.

### 1.7.1 Clean up an existing database

```
# Connect to the PostgreSQL database  
$ psql -U postgres  
  
# Drop the database  
tranalyzer=# DROP DATABASE tranalyzer;
```